

Lecture 2 - Solving inference problems using SMC



Thomas Schön

e-mail: thomas.schon@it.uu.se

Division of Systems and Control
Department of Information Technology
Uppsala University

Lecture 2 - Solving inference problems using SMC
Thomas Schön, Advanced Bayesian learning at Linköping University, Linköping, Sweden in May 2014.

Summary – Lecture 1 (I/III)

2(46)

A **model** is a compact and interpretable representation of the data that is observed. In this course, **model = PDF**!

Introduced the nonlinear state space model (SSM)

$$\begin{aligned}x_{t+1} | x_t &\sim f_{\theta,t}(x_{t+1} | x_t, u_t), \\ y_t | x_t &\sim g_{\theta,t}(y_t | x_t, u_t), \\ x_1 &\sim \mu_{\theta,t}(x_1), \quad (\theta \sim p(\theta)).\end{aligned}$$

Lecture 2 - Solving inference problems using SMC
Thomas Schön, Advanced Bayesian learning at Linköping University, Linköping, Sweden in May 2014.

Summary – Lecture 1 (II/III)

3(46)

State inference refers to the problem of finding information about the state(s) $x_{k:l}$ based on the available measurements $y_{1:t}$.

We showed the **key strategies** for state filtering and state smoothing. These strategies forms the foundation for what we will do today.

The state filtering problem: Compute the filtering pdf $p(x_t | y_{1:t})$.

Solved via the forward computations

$$\begin{aligned}p(x_t | y_{1:t}) &= \frac{\overbrace{g(y_t | x_t)}^{\text{measurement}} \overbrace{p(x_t | y_{1:t-1})}^{\text{prediction pdf}}}{p(y_t | y_{1:t-1})}, \\ p(x_t | y_{1:t-1}) &= \int \underbrace{f(x_t | x_{t-1})}_{\text{dynamics}} \underbrace{p(x_{t-1} | y_{1:t-1})}_{\text{filtering pdf}} dx_{t-1},\end{aligned}$$

Lecture 2 - Solving inference problems using SMC
Thomas Schön, Advanced Bayesian learning at Linköping University, Linköping, Sweden in May 2014.

Summary – Lecture 1 (III/III)

4(46)

We also derived the backward computations, where the **backward kernel** is key

$$p(x_t | x_{t+1}, y_{1:t}) = \frac{f(x_{t+1} | x_t) p(x_t | y_{1:t})}{p(x_{t+1} | y_{1:t})}.$$

Monte Carlo methods provide computational solutions to otherwise intractable problems relying on random sampling.

Importance sampling: Offers a way of evaluating integrals of the form

$$\int \varphi(z) \pi(z) dz.$$

Lecture 2 - Solving inference problems using SMC
Thomas Schön, Advanced Bayesian learning at Linköping University, Linköping, Sweden in May 2014.

1. Derive a first working particle filter (most common SMC sampler)
2. Particle filtering examples
 - a) An LGSS model
 - c) A nonlinear SSM
 - b) Indoor positioning
3. Targeting the JSD and particle degeneracy
4. Ancestor indices and an equivalent SMC formulation
5. Particle smoothers (very brief)
6. Some recent work

Recall that the nonlinear filtering problem amounts to computing the filter PDF $p(x_t | y_{1:t})$ when the model is given by

$$\begin{aligned} x_{t+1} | x_t &\sim f(x_{t+1} | x_t), \\ y_t | x_t &\sim g(y_t | x_t), \\ x_1 &\sim \mu(x_1). \end{aligned}$$

We have showed that the solution is

$$\begin{aligned} p(x_t | y_{1:t}) &= \frac{g(y_t | x_t)p(x_t | y_{1:t-1})}{p(y_t | y_{1:t-1})}, \\ p(x_t | y_{1:t-1}) &= \int f(x_t | x_{t-1})p(x_{t-1} | y_{1:t-1})dx_{t-1}. \end{aligned}$$

Relevant idea: Try to solve this using importance sampling!!

Algorithm 1 Importance sampler (IS)

1. Sample $z^i \sim q(z)$.
 2. Compute the weights $\tilde{w}^i = \tilde{\pi}(z^i) / q(z^i)$.
 3. Normalize the weights $w^i = \tilde{w}^i / \sum_{j=1}^N \tilde{w}^j$.
-


Each step is carried out for $i = 1, \dots, N$.

Resampling is the procedure that (randomly) **turns a weighted** set of samples $\{x_{t-1}^i, w_{t-1}^i\}_{i=1}^N$ **into an unweighted** set of samples $\{\tilde{x}_{t-1}^i, 1/N\}_{i=1}^N$ according to

$$\mathbb{P}(\tilde{x}_{t-1} = x_{t-1}^i | \{x_{t-1}^i, w_{t-1}^i\}_{i=1}^N) = w_{t-1}^i,$$

Resampling can be implemented in several ways and many algorithms exist.

Rather than going through all these algorithms, let us illustrate the meaning of resampling, resulting in one resampling algorithm.



Three new samples are generated in the figure above, corresponding to sample 2, 4 and 6.

Illustrating how resampling works (using 7 particles).

1. Compute the cumulative sum of the weights.
2. Generate $u \sim \mathcal{U}[0, 1]$.


Algorithm 2 Bootstrap particle filter

1. Initialization ($t = 1$):

- (a) Sample $x_1^i \sim \mu(x_1)$.
- (b) Compute the importance weights $\tilde{w}_1^i = g(y_1 | x_1^i)$ and normalize, $w_1^i = \tilde{w}_1^i / \sum_{j=1}^N \tilde{w}_1^j$.

2. for $t = 2$ to T do

- (a) Resample $\{x_{t-1}^i, w_{t-1}^i\}$ resulting in equally weighted particles $\{\tilde{x}_{t-1}^i, 1/N\}$.
- (b) Sample $x_t^i \sim f(x_t | \tilde{x}_{t-1}^i)$.
- (c) Compute the importance weights $\tilde{w}_t^i = g(y_t | x_t^i)$ and normalize $w_t^i = \tilde{w}_t^i / \sum_{j=1}^N \tilde{w}_t^j$.



The structure is the same for all PFs. For the bootstrap PF we have,

Resampling: $\{x_{t-1}^i, w_{t-1}^i\}_{i=1}^N \rightarrow \{\tilde{x}_{t-1}^i, 1/N\}_{i=1}^N$.

Propagation: $x_t^i \sim f(x_t | \tilde{x}_{t-1}^i)$.

Weighting: $\tilde{w}_t^i = W_t(x_t^i) = g(y_t | x_t^i)$ and normalize.

The result is a new weighted set of particles $\{x_t^i, w_t^i\}_{i=1}^N$.

“Whenever you are working on a nonlinear inference method, always make sure that it solves the linear special case first.”

Consider the following LGSS model (simple one dimensional positioning example)


$$\begin{pmatrix} p_{t+1} \\ v_{t+1} \\ a_{t+1} \end{pmatrix} = \begin{pmatrix} 1 & T_s & T_s^2/2 \\ 0 & 1 & T_s \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} p_t \\ v_t \\ a_t \end{pmatrix} + \begin{pmatrix} T_s^3/6 \\ T_s^2/2T_s \end{pmatrix} v_t, \quad v_t \sim \mathcal{N}(0, Q),$$

$$y_t = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} p_t \\ v_t \\ a_t \end{pmatrix} + e_t, \quad e_t \sim \mathcal{N}(0, R).$$

The KF provides the true filtering density, which implies that we can compare the PF to the truth in this case.

An LGSS example (II/II)

13(46)



Using 200 particles.

Using 20 000 particles.

The PF estimate converge as the number of particles tends to infinity.

Xiao-Li Hu, Thomas B. Schön and Lennart Ljung. **A Basic Convergence Result for Particle Filtering.** *IEEE Transactions on Signal Processing*, 56(4):1337-1348, April 2008.

D. Crisan and A. Doucet, **A survey of convergence results on particle filtering methods for practitioners.** *IEEE Transactions on Signal Processing*, vol. 50, no. 3, pp. 736-746, 2002.

Lecture 2 - Solving inference problems using SMC

Thomas Schön, Advanced Bayesian learning at Linköping University, Linköping, Sweden in May 2014.

A nonlinear example (I/II)

14(46)

Consider the following SSM (standard example in PF literature)

$$x_{t+1} = \frac{x_t}{2} + \frac{25x_t}{1+x_t^2} + 8 \cos(1.2t) + v_t, \quad v_t \sim \mathcal{N}(0, 0.5),$$

$$y_t = \frac{x_t^2}{20} + e_t, \quad e_t \sim \mathcal{N}(0, 0.5).$$

What is tricky with this model?


The best (only?) way of really understanding something is to implement it yourself.

Lecture 2 - Solving inference problems using SMC

Thomas Schön, Advanced Bayesian learning at Linköping University, Linköping, Sweden in May 2014.

A nonlinear example (II/II)

15(46)



True state (gray) and PF conditional mean estimate (black).

PF estimate of the filtering pdf $\hat{p}(x_{56} | y_{1:56})$.

Another indication that the conditional mean point estimate is dangerous.


Lecture 2 - Solving inference problems using SMC

Thomas Schön, Advanced Bayesian learning at Linköping University, Linköping, Sweden in May 2014.


Application example – indoor positioning (I/III)

16(46)

Aim: Compute the position of a person moving around indoors using sensors located in an ID badge and a map.




The sensors (IMU and radio) and the DSP are mounted inside an ID badge.




The inside of the ID badge.

Lecture 2 - Solving inference problems using SMC

Thomas Schön, Advanced Bayesian learning at Linköping University, Linköping, Sweden in May 2014.



pdf for an office environment, the bright areas are rooms and corridors (i.e., walkable space).



An estimated trajectory and the particle cloud visualized at a particular instance.

Show movies

This is work by Johan Kihlberg and Simon Tegelid in their MSc thesis entitled **Map aided indoor positioning**.

This work will also be presented at **Reglermötet** here in Linköping in the beginning of June.

1. Derive a first working particle filter (most common SMC sampler)
2. Particle filtering examples
 - a) An LGSS model
 - b) A nonlinear SSM
 - c) Indoor positioning
3. Targeting the JSD and particle degeneracy
4. Ancestor indices and an equivalent SMC formulation
5. Particle smoothers (very brief)
6. Some recent work

Our derivation of the PF is rather **non-standard**. The reason I like it is that it clearly shows why the resampling step is needed and where the need for the resampling step comes from.

The more **standard** way of deriving the PF is by **targeting the sequence of joint smoothing densities (JSD)** $\{p(x_{1:t} \mid y_{1:t})\}_{t \geq 1}$ (see lecture notes for details).

Algorithm 3 Sequential Monte Carlo (SMC)

1. Initialise by sampling $x_1^i \sim Q_1(x_1)$, setting $\tilde{w}_1^i = W_1(x_1^i)$, normalizing and setting $t = 1$.
2. **for** $t = 1$ **to** T **do**
 - (a) **Resample:** $\{x_{1:t}^i, w_t^i\}_{i=1}^N \rightarrow \{\tilde{x}_{1:t}^i, 1/N\}_{i=1}^N$.
 - (b) **Propagate:** propose new samples, $x_{t+1}^i \sim q_t(x_{t+1} | \tilde{x}_{1:t}^i)$ and set $x_{1:t+1}^i = \{x_{1:t}^i, \tilde{x}_{t+1}^i\}$.
 - (c) **Weight:** $\tilde{w}_{t+1}^i = W_t(x_{t+1}^i, \tilde{x}_{1:t}^i)$ and normalize.

There is a myriad of possible design choices available, some of them carrying their own name. The important thing is to understand the basic principles in action; the rest is just design choices.

Can you see any problems with the algorithm producing approximations of the JSDs according to


$$p(x_{1:t} | y_{1:t}) = \sum_{i=1}^N w_t^i \delta_{x_{1:t}^i}(x_{1:t})$$

The resampling step remove particles with small weights and duplicate particles with large weights.

This results in **particle degeneracy**, which we explain using a simple example.

Illustration of particle degeneracy (I/II)

Illustration of particle degeneracy (II/II)



Left plot: At each point in time all particles are plotted using a black dot and each particle is connected with its ancestor using a black line.

Right plot: The grey dots represents the $p(x_t | y_{1:t})$ at each point in time. The black lines shows the particle trajectories $\{x_{1:25}^i\}_{i=1}^{30}$ at time $t = 25$.

The right plot corresponds to the left plot with all trajectories that are not resampled removed (all particles are still visualized using gray dots).

This implies that if we are interested in the smoothing distribution

$$p(x_{1:T} \mid y_{1:T})$$

or some of its marginals we are **forced** to use different algorithms, which leads us to **particle smoothers**.

Resampling is the procedure that (randomly) **turns a weighted** set of samples $\{x_{t-1}^i, w_{t-1}^i\}_{i=1}^N$ **into an unweighted** set of samples $\{\tilde{x}_{t-1}^i, 1/N\}_{i=1}^N$ according to

$$\mathbb{P}(\tilde{x}_{t-1} = x_{t-1}^i \mid \{x_{t-1}^i, w_{t-1}^i\}_{i=1}^N) = w_{t-1}^i,$$

Using the **auxiliary variables** offered by the **ancestor indices** $\{a_t^i\}_{i=1}^N$ we can keep track of exactly what happens in each resampling step. This is important and valuable information.

Evolution of three particles for $t = 1, 2, 3$.




Figure: Evolution of a particle system. The ancestral path of x_3^1 , i.e. $x_{1:3}^1$, is shown as a thick line.

At time $t = 1$, particle x_1^2 is resampled twice and particle x_1^3 is resampled once. Hence, at time $t = 2$, the **ancestor indices** are

$$a_2^1 = 2, a_2^2 = 2 \text{ and } a_2^3 = 3.$$

Similarly, at time $t = 3$, the **ancestor indices** are given by

$$a_3^1 = 2, a_3^2 = 3 \text{ and } a_3^3 = 3.$$

The **ancestral path** of particle x_3^1 , which we denote $x_{1:3}^1$, is shown as a thick line in the figure. This path is given recursively from the ancestor indices,

$$x_{1:3}^1 = (x_1^{a_3^1}, x_2^{a_2^1}, x_3^1) = (x_1^2, x_2^2, x_3^1).$$

Introduce a sequence of proposal kernels on the product space $\{1, \dots, N\} \times \mathbf{X}$,

$$M_t(a_t, x_t) = w_{t-1}^{a_t} q_t(x_t | x_{1:t-1}^{a_t}).$$

Algorithm 4 Sequential Monte Carlo (SMC)

1. Initialise by sampling $x_1^i \sim q_1(x_1)$, setting $\tilde{w}_1^i = W_1(x_1^i)$, normalizing and setting $t = 1$.
 2. **for** $t = 2$ **to** T **do**
 - (a) Draw $\{a_t^i, x_t^i\} \sim M_t(a_t, x_t)$.
 - (b) Set $x_{1:t}^i = \{x_{1:t-1}^{a_t^i}, x_t^i\}$
 - (c) Compute $\tilde{w}_t^i = W_t(x_{1:t}^i)$ and normalize.
 3. **end**
-

The PF can be interpreted as a mechanism for generating a realization from the $2NT - N$ random variables $X_{1:T}, A_{2:T}$ described by the pdf

$$\psi(\mathbf{X}_{1:T}, \mathbf{A}_{2:T}) = \underbrace{\left\{ \prod_{i=1}^N q_1(X_1^i | y_1) \right\}}_{\text{Initialization}} \left\{ \prod_{t=2}^T \underbrace{\prod_{i=1}^N \mathcal{C}(A_t^i | \mathbf{w}_{t-1})}_{\text{Resampling}} \underbrace{q_t(X_t^i | x_{1:t-1}^{a_t^i}, y_{1:t})}_{\text{Propagation}} \right\}$$

which is defined on the space $\mathbf{X}^{TN} \times \{1, \dots, N\}^{(T-1)N}$.

The PF is a **combination of two sampling operations**.

The fact that the PF is a combination of two sampling operations can be taken one step further.

Combine the two sampling operations into one sampling operation using $M_t(a_t, x_t)$

$$\psi(\mathbf{X}_{1:T}, \mathbf{A}_{2:T}) = \left\{ \prod_{i=1}^N q_1(X_1^i | y_1) \right\} \prod_{t=2}^T \prod_{i=1}^N M_t(A_t^i, X_t^i | \mathbf{w}_{t-1}, x_{1:t-1}^{a_t^i}, y_{1:t})$$

1. Derive a first working particle filter (most common SMC sampler)
2. Particle filtering examples
 - a) An LGSS model
 - b) A nonlinear SSM
 - c) Indoor positioning
3. Targeting the JSD and particle degeneracy
4. Ancestor indices and an equivalent SMC formulation
- 5. Particle smoothers (very brief)**
6. Some recent work

By marginalizing

$$p(x_{1:T} | y_{1:T}) = p(x_T | y_{1:T}) \prod_{t=1}^{T-1} \frac{f(x_{t+1} | x_t) p(x_t | y_{1:t})}{p(x_{t+1} | y_{1:t})}.$$

w. r. t. $x_{1:t-1}$ and $x_{t+1:T}$ we obtain the following expression for the marginal smoothing pdf

$$p(x_t | y_{1:T}) = p(x_t | y_{1:t}) \int \frac{f(x_{t+1} | x_t) p(x_{t+1} | y_{1:T})}{p(x_{t+1} | y_{1:t})} dx_{t+1}.$$

Algorithm 5 Forward filtering/Backward smoothing (FFBSm)

1. Run the PF and store the particles and their weights $\{(w_t^i, x_t^i)\}_{i=1}^N$ for $t = 1, \dots, T$.
2. Initialise the smoothed weights $w_{T|T}^i = w_T^i$, for $i = 1, \dots, N$.
3. **for** $t = T - 1$ **to** 1 **do**
 - (a) Compute the smoothed weights

$$w_{t|T}^i = w_t^i \sum_{k=1}^N w_{t+1|T}^k \frac{f(x_{t+1}^k | x_t^i)}{v_t^k}, \quad v_t^k = \sum_{i=1}^N w_t^i f(x_{t+1}^k | x_t^i).$$

The resulting approximation $\hat{p}(x_t | y_{1:T}) = \sum_{i=1}^N w_{t|T}^i \delta_{x_t^i}(x_t)$ does not suffer from degeneracy.

$$p(x_{1:T} | y_{1:T}) = \left(\prod_{t=1}^T p(x_t | x_{t+1}, y_{1:t}) \right) p(x_T | y_{1:T}),$$

$$p(x_t | x_{t+1}, y_{1:t}) = \frac{f(x_{t+1} | x_t) p(x_t | y_{1:t})}{p(x_{t+1} | y_{1:t})}.$$

Strategy: Run a forward filter and generate samples backwards in time:

Sample $\tilde{x}_T \sim p(x_T | y_{1:T})$
 \vdots
 Sample $\tilde{x}_t \sim p(x_t | \tilde{x}_{t+1}, y_{1:t})$
 \vdots
 Sample $\tilde{x}_1 \sim p(x_1 | \tilde{x}_2, y_1)$

The resulting sample is then by construction $\tilde{x}_{1:T} \sim p(x_{1:T} | y_{1:T})$.

Key idea: Generate a backward trajectory by sampling from the following approximation

$$\hat{p}(x_t | \tilde{x}_{t+1}, y_{1:t}) = \sum_{i=1}^N \frac{w_t^i f(\tilde{x}_{t+1}^i | x_t^i)}{\underbrace{\sum_{k=1}^N w_t^k f(\tilde{x}_{t+1}^k | x_t^k)}_{\triangleq \tilde{w}_{t|T}^i}} \delta_{x_t^i}(x_t)$$

of the backward kernel. This amounts to making use of the backwards simulation strategy

$$\tilde{x}_T \sim \hat{p}^N(x_T | y_{1:T}),$$

$$\tilde{x}_t \sim \hat{p}^N(x_t | \tilde{x}_{t+1}, y_{1:t}),$$

to sample an entire backward trajectory $\tilde{x}_{1:T}$.

Still computationally very costly!

1. At time T draw among the particles according to

$$\mathbb{P}(\tilde{x}_T = x_T^i) = w_T^i.$$

2. At time $t = T - 1, \dots, 1$ draw among the particles according to

$$\mathbb{P}(\tilde{x}_t = x_t^i) = \tilde{w}_{t|T}^i,$$

Repeat M times, resulting in

$$\hat{p}^M(x_{1:T} | y_{1:T}) = \frac{1}{M} \sum_{j=1}^M \delta_{\tilde{x}_{1:T}^j}(x_{1:T})$$

Algorithm 6 Particle based FFBSi

1. Run a PF to generate $\{x_t^i, w_t^i\}_{i=1}^N$ for $t = 1, \dots, T$.
2. **Initialization** ($t = T$):
 - (a) Sample independently $\{b_T(j)\}_{j=1}^M \sim \mathcal{C}(\{w_T^i\}_{i=1}^N)$.
 - (b) Set $\tilde{x}_T^j = x_T^{b_T(j)}$ for $j = 1, \dots, M$.
3. **for** $t = T - 1$ **to** 1 **do**
4. **for** $j = 1$ **to** M **do**
 - (a) Compute $\tilde{w}_{t|T}^{ij} \propto w_t^i f(\tilde{x}_{t+1}^j | x_t^i)$ and normalize.
 - (b) Sample independently $b_t(j) \sim \mathcal{C}(\{\tilde{w}_{t|T}^{ij}\}_{i=1}^N)$
 - (c) Set $\tilde{x}_t^j = x_t^{b_t(j)}$ and $\tilde{x}_{1:T}^j = \{\tilde{x}_t^j, \tilde{x}_{t+1:T}^j\}$.


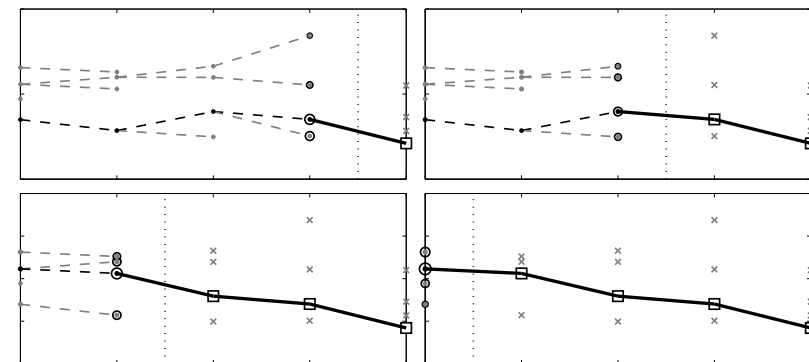


Illustration of a particle system $\{x_t^i, w_t^i\}_{i=1}^4$ generated using $N = 4$ particles during $T = 5$ time steps. The size of the particle visualize the corresponding weight.



The backward smoothing weights $\tilde{w}_{t|T}^i$ are visualized using circles.


The BSi samples a trajectory from the empirical joint smoothing density by first choosing a particle x_T^m with probability proportional to w_T^m and then generating the **ancestral path** $b_{1:T}$ according to the backward kernel.

The resulting trajectory is

$$x_{1:T}^* = \{x_1^{b_1}, \dots, x_T^{b_T}\}.$$


The computational cost is $\mathcal{O}(N^2)$, which is too expensive to be practical.

Recall that the **main problem** with using the PF to sample from $p(x_{1:T} | y_{1:T})$ is the particle degeneracy problem.



In the figures we show two samples from $p(x_{1:T} | y_{1:T})$ generated using the PF. The diversity is completely gone from time ≈ 30 and backwards.

The backward simulator **resolves the particle degeneracy problem**.



Note that in the above samples there is no loss of diversity.

To obtain a practical algorithm we have to reduce the computational complexity, how?!

Key insight: We do not have to compute all the smoothing weights $\{\tilde{w}_{t|T}^{ij}\}$ to be able to sample from the empirical backward kernel.

How: Use rejection sampling to **sample the indices**! This means that we should not have to compute N indices just to draw one sample.

For details, see Chapter 8 and/or

Frederik Lindsten and Thomas B. Schön. **Backward simulation methods for Monte Carlo statistical inference**. *Foundations and Trends in Machine Learning*, 6(1):1-143, 2013.

New result: We have been able to construct and learn a Gaussian process (GP) state space model

$$\begin{aligned} f(x_t) &\sim \mathcal{GP}(m_{\theta_x}(x_t), k_{\theta_x}(x_t, x'_t)), \\ x_{t+1} | f_t &\sim \mathcal{N}(x_{t+1} | f_t, Q), \\ y_t | x_t &\sim p(y_t | x_t, \theta_y). \end{aligned}$$

Key idea: Marginalize out the function f .

Problem: Renders the model non-Markovian. **Solution:** PG-AS

Roger Frigola, Fredrik Lindsten, Thomas B. Schön and Carl E. Rasmussen. **Bayesian inference and learning in Gaussian process state-space models with particle MCMC**. In *Advances in Neural Information Processing Systems (NIPS) 26*, Lake Tahoe, NV, USA, December 2013.

Roger Frigola, Fredrik Lindsten, Thomas B. Schön and Carl E. Rasmussen. **Identification of Gaussian Process state-space models with particle stochastic approximation EM**. In *Proceedings of the 18th World Congress of the International Federation of Automatic Control (IFAC)*, Cape Town, South Africa, August 2014. (accepted for publication).

SMC samplers are used to approximate a sequence of probability distributions on a sequence of probability spaces.

Constructing an artificial sequence of intermediate target distributions for an SMC sampler is a powerful (and **quite possibly underutilized**) idea.

Key idea: Perform and make use of a **sequential decomposition** of the graphical model.

See talk at the Newton institute two weeks ago:

<http://www.newton.ac.uk/programmes/MCM/seminars/2014042510401.html>